QUIZIZZ

Cloud Architecting - Week 3
20 Questions

NAME : Ng Yin Xuan

CLASS : RSW3S1G3

DATE : 7/7/2025

1. Which component of the AWS Global Infrastructure does Amazon Cloudfront use to ensure low-latency delivery?

☐ AWS Regions   ⬤ AWS edge locations

☐ AWS Availability Zones   ☐ Amazon Virtual Private Cloud (Amazon VPC)

2. You can run applications and workloads from a Region closer to the end users to _____ latency

☐ increase   ⬤ decrease

3. Networking, storage, compute, and databases are examples of service categories that AWS offers.

☐ False   ⬤ True

4. Which of the following are geographic areas that host two or more Availability Zones?

☐ AWS Origins   ⬤ AWS Regions

☐ Compute zones   ☐ Edge locations

5. ____ means the infrastructure has built-in component redundancy and ___ means that resources dynamically adjust to increases or decreases in capacity requirements.

⬤ Fault tolerant, elastic and scalable   ☐ Elastic and scalable, no human intervention

☐ Fault tolerant, no human intervention   ☐ No human intervention, fault tolerant

6. Availability Zones within a Region are connected through low-latency links.

⬤ True   ☐ False

7. Which of these statements about Availability Zones is not true?

☐ Availability Zones are made up of one or more data centers.

☐ Availability Zones are connected to each other using high-speed private links.

☐ Availability Zones are designed for fault isolation.

⬤ A data center can be used for more than one Availability Zone.

8. What is true about Regions? (Choose two.)

☐ They are the physical locations of your customers.

☐ All Regions are located in one specific geographic area.

⬤ Each Region is located in a separate geographic area.

⬤ A Region is a physical location that has multiple Availability Zones.

9. AWS highly recommends provisioning your compute resources across _____ Availability Zones.

⬤ multiple

☐ single

☐ no

☐ all

10. Edge locations are only located in the same general area as Regions.

⬤ False

☐ True

11. Amazon Simple Storage Service (Amazon S3) provide a good solution for which of the following use cases?

☐ A cluster for traditional Apache Spark and Apache Hadoop installations to process big data

⬤ An internet-accessible storage location for video files that an external website accesses

☐ A data warehouse for business intelligence

☐ Hourly storage of frequently accessed temporary files

12. A company is interested in using Amazon Simple Storage Service (Amazon S3) alone to host their website, instead of a traditional web server. Which types of content does Amazon S3 support for static web hosting? (Select THREE.)

⬤ Video and sound files

⬤ Client-side scripts

☐ Server-side scripts

⬤ HTML files and image files

13. Which scenarios represent a good use for Amazon Simple Storage Service (Amazon S3)? (Select TWO.)

- ● Backing up critical data
- ● Storing computation and analytics data
- ☐ Housing the root volume of a live operating system
- ☐ Exposing a virtual tape library to on-premises backup systems

14. A company wants to use an S3 bucket to store sensitive data. Which actions can they take to protect their data? (Select TWO.)

- ☐ Enabling server-side encryption on the S3 bucket after uploading sensitive data
- ☐ Using Secure File Transfer Protocol (SFTP) to connect directly to Amazon S3
- ● Enabling server-side encryption on the S3 bucket before uploading sensitive data
- ● Using client-side encryption to protect data in transit

15. A company must create a common place to store shared files. Which requirements does Amazon Simple Storage Service (Amazon S3) support? (Select TWO.)

- ☐ Attach comments to files
- ● Recover deleted files
- ● Maintain different versions of files
- ☐ Lock a file so that only one person at a time can edit it

16. A customer service team accesses case data daily for up to 30 days. Cases can be reopened and require immediate access for 1 year after they are closed. Reopened cases require 2 days to process. Which solution meets the requirements and is the most cost-efficient?

- ● Store case data in S3 Standard. Use a lifecycle policy to move the data into S3 Standard-Infrequent Access (S3 Standard-IA) after 30 days.
- ☐ Store all case data in S3 standard so that it is available whenever needed.
- ☐ Store case data in S3 Intelligent-Tiering to automatically move data between tiers based on access frequency.
- ☐ Store case data in S3 Standard. Use a lifecycle policy to move the data into Amazon S3 Glacier after 30 days.

17. Which Amazon Simple Storage Service (Amazon S3) unaccelerated data transfers have an associated cost? (Select TWO.)

- ☐ IN from the internet
- ☐ OUT to other AWS services in the same Region
- ● OUT to other AWS Regions
- ● OUT to the internet

18. A company is migrating 100 terabytes (TB) of data from their on-premises data centre to Amazon S3. The company connect AWS by using a single 155 Mbps internet connection. Which data transfer option is the fastest and most cost-effective?

- ● AWS Snowball
- ☐ Amazon Management Console
- ☐ Amazon S3 multipart uploads
- ☐ AWS Snowmobile

19. A video producer must regularly transfer several video files to Amazon S3. The files range from 100-700 MB. The internet connection has been unreliable, causing some uploads to fail. Which solution provides the fastest, most reliable, and most cost-effective way to transfer these files to Amazon S3?

- ☐ AWS Snowball
- ☐ AWS Snowmobile
- ● Amazon S3 multipart uploads
- ☐ Amazon Management Console

20. Which qualities vary by AWS Region? (Select TWO.)

- ☐ High availability of workloads
- ● Service and feature availability
- ☐ Data privacy
- ● Cost-effectiveness of workload